

Two Regimes of Chain-of-Thought Unfaithfulness: Metric-Based Detection Fails Where Models Are Wrong

Anonymous ACL submission

Abstract

Detecting unfaithful chain-of-thought (CoT) reasoning is usually evaluated in aggregate, over traces the deployed model got right and wrong alike. Stratifying the only expert-annotated instance-level benchmark by answer correctness splits detection into two regimes with opposite characters. On incorrect answers—which carry 69% of annotated unfaithfulness—no metric-based signal is detectably above chance; on correct answers, several are moderately informative and the standard step-removal metric is confidently *inverted*, anti-correlating with human labels (0.349 full-set AUROC), an inversion that reproduces on the benchmark’s own released scores and under intervention-defined labels. The decomposition is exact: aggregate metric scores are dominated by cross-regime comparisons that correctness alone resolves, so answer incorrectness—an oracle—matches or beats every purpose-built metric, and rankings reorder under composition changes with per-pair behavior frozen. A prompted LLM judge is the exception: it beats the correctness oracle and keeps significant blind-regime signal, but degrades exactly there ($0.830 \rightarrow 0.679$), robustly across prompts and judge models; ablations show it reads the trace, not the question. Constructed training data does not close the gap: instructed rationalization is internally detectable in all seven models tested yet transfers to neither annotated regime, and the one apparent hint-construction transfer attenuates on matched question cohorts. A pre-registered independent test then *fails*: on an intervention-labeled dataset with disjoint generators, both hypotheses are rejected—the judge falls significantly below chance (0.419; 0.332 with generator composition controlled) while step count reaches 0.878. Two operationalizations of instance-level unfaithfulness order our detectors oppositely, so we scope the ordering to expert annotation and leave the reversal unexplained. We draw prescriptions for evaluating

and validating faithfulness detectors.

1 Introduction

Large language models produce fluent step-by-step explanations that appear to expose a reasoning process, but the stated reasoning is sometimes not the process that produced the answer (Lanham et al., 2023; Turpin et al., 2023; Chen et al., 2025). As CoT becomes an oversight surface, where an auditor reads the reasoning to decide whether to trust the answer, the practical question is instance-level: can *this* explanation be trusted? FaithCoT-Bench (Shen et al., 2026b) formalizes this problem with human annotations, evaluating eleven detectors and finding metric-style families weak while LLM-as-judge methods are strongest in aggregate.

Detection is ordinarily evaluated in *aggregate*, pooling traces the deployed model answered correctly and incorrectly. We audit that practice on the only expert-annotated instance-level benchmark and find it conceals two regimes with opposite characters: on incorrect answers—69% of annotated unfaithfulness—no metric-based signal is detectably above chance, while on correct answers several are moderately informative and the standard step-removal metric is confidently *inverted*. The decomposition is exact: aggregate AUROC is dominated by cross-regime comparisons that answer correctness alone resolves, so an incorrectness oracle matches or beats every purpose-built metric, and detector rankings reorder under composition alone. A prompted judge partially escapes—beating the oracle, retaining degraded blind-regime signal, reading the trace rather than the question—while constructed training data does not close the gap: instructed rationalization transfers to neither regime, and the one apparent hint transfer attenuates on matched question cohorts. Figure 1 summarizes; Figure 3 shows the phenomenon itself: the same model, on the same problem, fails cleanly; handed a hint, it produces a lo-

cally flawless derivation of the hinted answer without mentioning the hint. Nothing on the surface distinguishes that trace from genuine reasoning.

Contributions: (1) a correctness-stratified evaluation protocol for instance-level faithfulness detection—prior evidence of faithfulness—accuracy coupling is model-level (Bentham et al., 2024) or unstratified—with the two-regime finding, an exact composition decomposition of aggregate scores, and the demonstration that detector rankings reorder under composition alone; (2) an answer–reasoning coupling *account* and regime localization for the step-removal inversion, complementing constructed-ground-truth metric failures (Gur-Arieh et al., 2026; Zaman and Srivastava, 2025a) with human-label evidence; (3) a regime-stratified judge and internals analysis: the judge survives the blind regime degraded, reads the trace rather than the question, and does not subsume linear probes; (4) a controlled *transfer-validation of constructed unfaithfulness against human annotations*: instructed rationalization (cf. Kirch et al., 2025) is internally detectable in every model tested yet carries no annotated-regime signal (bounded below 0.55 in one cell), and the single apparent hint transfer attenuates on matched cohorts, sensitive to training-question selection in both classes.

We proceed in three parts. Part I (§4) asks what behavioral detection actually measures, and finds metric-based signals largely re-detect answer correctness while the standard metric is inverted. Part II (§5–§5.1) shows the failure is structured rather than uniform: stratifying by correctness yields two regimes with opposite behavioral profiles and different internal decodability. Part III (§6) then asks whether the constructions used to build detector data share representations with the annotated phenomenon at all—the paper’s central test.

2 Related Work

CoT faithfulness. Lanham et al. (2023) introduced perturbation-based faithfulness tests, from which step-removal “answer-tracing” metrics descend. Turpin et al. (2023) showed models rationalize biased answers without verbalizing the bias; Chen et al. (2025) measured low hint-verbalization rates; Arcuschin et al. (2025) document post-hoc rationalization in the wild; Zaman and Srivastava (2025b) show unverbali-

zed hints still causally flow through the CoT—our hint-induced construction operationalizes that paradigm as a labeled testbed. Bentham et al. (2024) showed model-level unfaithfulness scores largely track accuracy; we give the instance-level counterpart against human labels. FaithCoT-Bench (Shen et al., 2026b) contributes the annotations we evaluate against and notes that correctness and faithfulness diverge at the instance level; we quantify the residual coupling, control for it, and resolve a label-semantics mismatch in its documentation (§3).

Concurrent work (2026). Reading unfaithfulness from internals is active: Mirtaheri and Belkin (2026) find post-CoT activation probes beat CoT monitors on hint-induced motivated reasoning; Occhipinti et al. (2026) probe and steer unfaithful spans; Shen et al. (2026a) (the FaithCoT-Bench group) build a circuit-based detector; Gur-Arieh et al. (2026) report below-chance AUROCs for perturbation metrics against constructed ground truth; Cox et al. (2026) decode answers pre-CoT. None stratifies *human-annotated* detection by correctness, explains the metric inversion via that structure, or compares instructed, hint-induced and annotated unfaithfulness at the representation level—the three additions here. Kirch et al. (2025) show probe performance degrades under off-policy and distribution shift; our transfer analysis extends that caution to CoT faithfulness.

Structural verification. PARC (Mukherjee et al., 2025), GoV (Fang et al., 2025) and VeriCoT (Feng et al., 2025) verify chains structurally, targeting correctness and logical validity; we find such structure largely does not transfer to *faithfulness*. Step-level benchmarks (Zheng et al., 2024; Jacovi et al., 2024; Pham et al., 2026) evaluate related but distinct constructs.

Probing. Our probing method is deliberately standard (linear probes with selection-corrected permutation testing): the contribution is the behavioral-vs-internal *contrast* and the transfer analyses, not a new technique. Following Belinkov (2022), decodability does not imply the model *uses* the information (Elazar et al., 2021); probe claims are therefore stated as “linearly decodable,” with steering as separate (weak) causal evidence.




Figure 1: Overview. Stratifying FaithCoT-Bench by correctness splits detection into two regimes, behaviorally and internally; of the constructed testbeds, only hint-induced flip probes transfer to the annotations, onto the metric-blind regime.

3 Setup

Benchmark. FaithCoT-Bench spans four domains (LogiQA (Liu et al., 2020), TruthfulQA (Lin et al., 2022), AQUA (Ling et al., 2017), HLE-Bio (Phan et al., 2025)) and four models (Llama-3.1-8B-Instruct, Qwen-2.5-7B-Instruct, GPT-4o-mini, Gemini-2.5-Flash): 1,364 released traces, 1,304 carrying a binary human label and 1,303 also a four-way code (census in Appendix B). It supplies derived metrics—notably *soft_faithfulness*, a step-removal answer-tracing score—and *human* annotations: binary unfaithfulness and four-way *faithful_type*. We evaluate against the human labels throughout.

Label semantics, verified against the data. At the time of our analysis the benchmark’s documentation and its released data disagreed about what the four *faithful_type* codes mean. In the data they pair as **ft1 faithful / ft2 unfaithful on incorrect answers; ft3 faithful / ft4 unfaithful (post-hoc) on correct answers**; the then-current documentation stated the reverse. The faithful/unfaithful axis is anchored by the separate binary label field (95.7% co-occurrence; same annotation effort, stored separately), and we resolved the correctness axis three independent ways: cross-tabulating the release’s *own stored* parsed answers against its gold labels (near-deterministic in every domain), reproducing the benchmark paper’s per-model accuracy statistics (which match only under the data-side semantics), and checking released per-trace metric scores (Appendix B). The discrepancy was independently flagged on the repository’s public issue tracker and corrected

upstream¹; we use the verified data-side semantics. Counts: ft1 281, ft2 233, ft3 682, ft4 107; hence **233/340 ≈ 69% of annotated unfaithfulness sits on incorrect answers**. The mismatch concerned documentation only, not the labels or the benchmark’s reported results. The methodological point—verify label semantics against raw data, since any analysis inheriting the documented pairing silently swaps its regimes—is the first of the three prescriptions in §7.

Statistical standards. Detection claims use AU-ROC with 2,000-resample percentile-bootstrap 95% CIs (minority classes 13–43%); because traces cluster by model and domain, headline effects are additionally checked with a cluster bootstrap over the 8 model×domain cells (both survive: incorrectness [0.616, 0.753]; inverted answer-tracing [0.575, 0.728]). Probe claims use held-out splits plus *selection-corrected* permutation tests: because the best layer is chosen post hoc, the null takes the max (or mean, matching the reported statistic) over layers under label permutation, coupled across layers (Nichols and Holmes, 2002):

$$T = \max_{\ell \in L} \text{AUROC}_{\ell} \text{ (or its mean over } L),$$

$$p = \frac{1 + \sum_{b=1}^B \mathbf{1}[T^{(b)} \geq T^{\text{obs}}]}{B + 1}. \quad (1)$$

Values reported as 0.005 mean $p \leq 1/201$; headline transfer tests use 1,000 permutations, and the null doubles as the memorization control (Hewitt and Liang, 2019). Pooled judge estimates re-

¹Independently flagged on the repository’s public issue tracker in July 2026; the documentation was corrected upstream later that month. Direct links and exact dates withheld for anonymity; included in the camera-ready.

computed with a question-clustered bootstrap (340 clusters) are essentially unchanged; trace-level intervals are reported throughout. Δ AUROC comparisons are descriptive, not a confirmatory family. Scalers and PCA are fit within training folds only.

Reproducibility and artifacts. An anonymized supplement contains the scripts and per-experiment result files backing every reported number, with a claims-to-scripts map; checkpoints and generation settings are in Appendix H. The constructed testbeds (898 hint-flip, 2,379 genuine traces) will be released with the camera-ready. Experiments ran on two RTX 3070 (8 GB) GPUs with 4-bit NF4 quantization (Dettmers et al., 2023).

4 Part I: What Behavioral Detection Measures

Signal families. *Answer-tracing:* sensitivity of the final-answer distribution to removing individual steps; the benchmark’s step-removal metric is, for a trace with T steps and stored parsed answer a^* ,

$$\text{soft_faithfulness}(x) = \frac{1}{T} \sum_{i=1}^T |\delta_i|, \quad (2)$$

where $\delta_i = P(a^* \mid \text{CoT}) - P(a^* \mid \text{CoT} \setminus \text{step}_i)$ is the shift in the answer’s probability when step i is deleted; higher means the answer depends on more of its stated steps (we use the benchmark’s stored scores). *Prefix instability* (PI) is an answer-tracing variant: the mean L_1 change between successive prefixes’ answer distributions. *Premise-DAG structure* (PARC-style graphs over steps) and *graph-targeted interventions* (corrupting the most load-bearing premise; LLM extractor validated at recall 0.82 / F1 0.79 vs. PARC gold) form the structural families, and *per-step NLI support* scores entailment of each step from question and prior steps (RoBERTa-large-MNLI; Liu et al., 2019; Williams et al., 2018); definitions in Appendix H. Because behavioral monitoring in practice often means *prompting an LLM judge*, we add one as a fifth family: GPT-4o-mini (temperature 0, one fixed prompt written before seeing any output, JSON-forced 0–100 score) over all 1,303 labeled traces, shown the question, options and full CoT but never the gold answer. That judge methods lead *in aggregate* is the benchmark’s own finding (Shen et al., 2026b); what it does not report is

judge performance stratified by correctness. Signal directions are fixed a priori; where a signal discriminates only in the *opposite* direction we report it as inverted—one sign selection made on the full data and reused unchanged in the stratified analyses.

Results. Figure 2 (left; full values in Table 2): **answer incorrectness alone has the highest AUROC among the metric-based signals (0.696)**—an oracle diagnostic, since it requires the gold answer. Paired oracle–metric Δ AUROC: +0.071 (PI) and +0.128 (NLI), both resolved; +0.045 (answer-tracing, not resolved); CIs in Table 2. Annotated unfaithfulness concentrates on incorrect answers, so metric-based detection largely reduces to accuracy prediction. **The prompted judge is the only behavioral signal to exceed the oracle:** judge 0.782 [0.756, 0.808] vs. oracle 0.703, paired Δ +0.079 [+0.043, +0.117] (complete-feature subset 0.771 vs. 0.697). Behavioral detection is therefore not *only* accuracy prediction, though the judge partially tracks incorrectness (0.684). Structural families add nothing significant (DAG max lookback is weakly informative on AQUA, 0.636). The 633-trace subset is set by feature availability, not outcomes.

Metric inversion. `soft_faithfulness` used as intended scores **0.349 [0.305, 0.389]**, significantly *below* chance; unfaithful traces score *higher* by +0.139 [+0.096, +0.182]. This is not a pipeline artifact: recomputing from the benchmark’s *own released* per-trace scores against its own human labels reproduces it (0.348; per-model values in Appendix B). Concurrent meta-evaluations report below-chance AUROCs against constructed ground truth (Gur-Arieh et al., 2026) and failures to beat random baselines (Zaman and Srivastava, 2025a); ours is against *human* labels, and the regime decomposition supplies a mechanism. Mean scores by annotation: ft3 (faithful-correct) **0.230** vs. ft1 0.371, ft2 0.397, ft4 **0.405**—the metric scores faithful-correct traces *least* faithful, plausibly because a model that genuinely solved the problem can re-derive its answer when steps are removed, while a tightly answer-coupled rationalization cannot. Within the incorrect regime the metric is merely uninformative (0.470, n.s.); within the correct regime it is confidently inverted (0.333). On these data it tracks answer–reasoning *coupling* rather than annotated faithfulness (§7), and §6 replicates the inversion on causally-labeled

traces.

5 Part II: Two Detection Regimes

Stratifying the complete-feature subset by answer correctness gives each half its own detection problem, with opposite characters (Figure 2; exact values in Table 5, Appendix B). The asymmetry is unchanged when the separate binary unfaithfulness field replaces `faithful_type` as the target (Appendix B). Per-cell heterogeneity is real (Appendix I): both headline effects are near chance on HLE-Bio ($n=40$ per model).

The detectable regime (correct answers). Separating faithful reasoning (ft3) from post-hoc rationalization (ft4) on correct answers is *partially* solvable behaviorally: answer-tracing (inverted), prefix instability (inverted), and NLI step-support all clear chance (0.63–0.67). The inversion is sharpest here: intended-direction `soft_faithfulness` scores 0.333, *worse than a coin flip* on exactly the distinction the metric was designed for.

The metric-blind regime (incorrect answers). Distinguishing honest from unfaithful error is opaque to every metric-based signal: **we detect no above-chance performance from any of the four algorithmic families**; every CI includes 0.5 (intervals are wide, so modest effects cannot be excluded). A *learned* text baseline also fails (0.446 / 0.474), and the pattern replicates on the full four-model labeled set (incorrect regime all n.s.; correct regime NLI 0.622). That benchmark-wide carrier is largely *length*—step count alone reproduces the asymmetry, while answer-tracing and PI survive length-partialing (Appendix I). Correct-regime detection is thus moderate but partly *shallow*; in the incorrect regime not even length carries signal. The judge is the exception: it retains significant signal (0.679 [0.633, 0.721]) but degrades sharply from 0.830 (Δ 0.152 [0.089, 0.211], $p<0.001$; robust under both decompositions, Appendix I). Neither our prompt wording nor swapping in a stronger judge changes this: an independently worded prompt reproduces the degradation (0.786 $\rightarrow$ 0.667), as does gpt-4o (0.786 $\rightarrow$ 0.670), which does *not* rescue the blind regime and mostly buys correctness tracking (0.745 vs. 0.684)—a two-judge comparison, not a scaling law. Ablations locate the signal: with the CoT withheld the judge is at or below chance ev-

erywhere but one marginal cell (full 0.487; correct 0.552); on the CoT alone it keeps most of its signal (0.693; blind regime 0.661) while losing most of its correctness tracking (0.560 vs. 0.684)—it reads the trace, not the question’s difficulty. Within a regime correctness is constant, so this is not correctness re-detection. Regime-difference tests: answer-tracing +0.136 [+0.035, +0.235], PI +0.178 [+0.078, +0.279], NLI +0.085 (n.s.), judge +0.152 [+0.089, +0.211].

An independent test, pre-registered and failed.

We registered endpoints, population, margin and analysis before computing any value against a second dataset: BonaFide (Gur-Arieh et al., 2026), whose whole-trace labels use construction-specific process/source-claim criteria rather than the same expert annotation standard and whose ten generators do not overlap ours. On its incorrect-answer population ($n=1,113$; 945/168) **both hypotheses failed**: the judge is significantly *below* chance (0.419 [0.372, 0.464]; gpt-4o 0.270) and the metric families are not null (step count 0.878, NLI 0.826). Conditioning separates them: feature signals decay toward chance (0.521 / 0.603 with generator, task and length fixed) while the judge inversion *strengthens* (0.193). That asymmetry is the finding: the inversion survives every restriction applied (instruct-only 0.332; four of five generators separately). **We scope the ordering above to expert-annotated unfaithfulness: it does not transfer here.** A separate exploratory GRACE check (Pham et al., 2026), on human step labels with the tested NLI procedures, likewise leaves every correctness-stratum interval covering zero, while an answer-error association persists under several aggregation rules; assistant-derived equivalence reviews and wide intervals limit it (Appendix K). Why the judge inverts is unresolved: judged scores are higher for short honest errors than for elaborate unacknowledged-hint rationalizations, consistent both with the coupling account of §7.2 and with a rubric–construct mismatch; we do not adjudicate between them (Appendix J).

5.1 What internals add

We fit standardized PCA-50 logistic-regression probes on hidden states at every layer, read at the final sequence position, with 5-fold CV and the selection-corrected test of §3 (teacher-forced re-encoding; the two open models). **Each regime is decodable—in a different model.** On ft1 vs.




Figure 2: The audit in one view (complete-feature subset; metric families). Left: incorrectness alone (red) matches or exceeds every purpose-built metric signal. Middle: on *incorrect*-answer traces no metric signal is detectably above chance (dashed line). Right: on *correct*-answer traces three metric signals are moderately informative (blue = CI excludes 0.5). Bars: bootstrap 95% CIs. The prompted LLM judge (gold, black edge) is evaluated on the *same* subset (0.771 / 0.635 [0.568, 0.702] / 0.832 across the panels; full labeled set: 0.782 / 0.679 / 0.830).

ft2, where every metric-based signal fails, Llama-3.1-8B reaches held-out AUROC 0.67 ± 0.08 under nested layer selection, generalizing across domains (LODO 0.60–0.71); Qwen is weak and not significant—“not detected,” not “absent” ($n=126$; cf. §6.3 (3)). On ft3 vs. ft4 the pattern *flips*: Qwen is significant (0.701, $p=0.014$) while Llama’s higher curve (0.759) does not survive selection correction ($p=0.19$; ft4 $n=26$), underpowered rather than negative. Probes trained on one regime do not transfer to the other in either model or direction (layer-means 0.52–0.57, all n.s.; Table 3). With the flipped model-dependence, this is most consistent with unfaithful *error* and post-hoc *rationalization* being distinct phenomena on different directions—a hypothesis, not an established mechanism; an exploratory steering check points the right way (0.22 vs. 0.08 at $+6\sigma$; Appendix C).

6 Part III: Do Constructions Measure the Real Phenomenon?

The annotations supply only 221 unfaithful traces across both regimes in the two probed models (340 benchmark-wide). Scaling requires *constructing* unfaithfulness—and the construction itself turns out to be the object of study.

6.1 Instructed construction: 7 models

For math problems (AQuA (Ling et al., 2017); GSM8K (Cobbe et al., 2021)) with gold answer a^* , the same model either solves step-by-step (kept only when correct: *genuine*) or is instructed to justify the given a^* (*post-hoc*). Both end at the same correct answer; probed contexts are identical (the instruction is never probed) and decoding is matched (Appendix E). Across **seven open models spanning three families and 2024–2026** (Grattafiori et al., 2024; Qwen Team et al., 2024; Gemma Team et al., 2024; DeepSeek-AI et al., 2025), the construction is permutation-significant in *every* model (Figure 6, Table 7; Appendix F), strongest in standard instruct models and weakest in reasoning-distilled ones. Surface-feature base-lines stay near chance for five of seven; Llama (0.64) and Gemma-4 (0.67) carry partial surface tells that the probes clear by ≥ 0.10 . Qwen-2.5 is strongest here, consistent with its regime pattern (§5.1): like the correct-answer regime, this construction pairs traces that all end at correct answers.

6.2 Hint-induced flip construction

Following the biased-context paradigm (Turpin et al., 2023; Chen et al., 2025), we first ask each problem clean (correct $\Rightarrow$ *genuine*, capped). For problems the model got *wrong*, we re-ask with a

sycophancy hint toward the gold answer (“I think the answer is a^* ...”; cf. Sharma et al., 2024). A trace is kept as *hint-induced* iff the model flips to a^* and the CoT never references the hint (mention-filtered; rejected cases are honest deference). The label is **intervention-defined at the sample level**, satisfying the unfaithfulness criterion of Chen et al. (2025). Decoding is sampled ($T=0.7$) and clean/hinted randomness is not coupled per item, so the label records an *observed* conditional flip, not a within-pair counterfactual; it is not a problem-level inability claim (one independent resample solves 59% of these problems; only 25%/29% fail two consecutive resamples—Appendix G). We distinguish the established claim (“the sampled answer differed under the hint, without verbalization”) from the stronger reading (“the CoT was generated answer-first”), which remains an inference. Probed contexts exclude the hint text. On the full AQUA-RAT and GSM8K test sets this yields **185 hint-induced + 428 genuine traces for Llama** and **70 + 464 for Qwen**, with an exhaustive keyword audit confirming no leakage. One caveat: genuine and hint-induced traces come from *disjoint problem pools*, so within-testbed decodability may partly reflect difficulty; the transfer analysis below is the difficulty-robust claim.

Both models are strongly decodable on this set (held-out 0.752 / 0.835, $p \leq 0.001$), replicated under a second, impersonal template (0.731 / 0.824) and on two further source domains built identically and matched on option count, LogiQA (0.672 / 0.793) and TruthfulQA (0.726 / 0.808; Appendix F). Input and text baselines are substantial here and we report them prominently—surface 0.626/0.703, learned text 0.611/0.751, question-only 0.646/0.664—so within-testbed decodability, especially for Qwen, is partly difficulty- and text-explainable, the CoT probes exceeding the strongest such baseline by +0.08–0.11. Most strikingly, **the metric inversion replicates under intervention-defined labels**: intended-direction soft_faithfulness scores **0.389 [0.301, 0.481] (Llama)** and **0.251 [0.166, 0.350] (Qwen)** on the AQUA portion where it is computable—on traces selected by an observed wrong→hinted flip, the standard metric confidently points the wrong way. This confirms the *coupling* account under coupling-aligned labels: genuine traces come from baseline-solvable problems by construction, so low deletion-sensitivity is expected there.

6.3 Transfer against the annotated regimes

Do the constructions share the annotated regimes’ representations? We train probes on each distribution and test on the others, sweeping all layers (Table 1; heatmap view in Appendix D, Figure 5). We make the following observations.

(1) Instructed rationalization does not transfer, and the null is bounded. Despite being cleanly detectable within-distribution in all seven models, its probes sit at or below chance against the incorrect regime (layer-mean 0.431). For one cell this is stronger than a failure to reject: the bootstrap upper bound is 0.502, excluding a 0.55 margin (the Qwen and correct-regime cells are unbounded nulls). It is not an artefact of linearity—a source-selected nonlinear head, gated to match the linear probe in-distribution, leaves it intact ($0.431 \rightarrow 0.474$)—nor of weak probes: on the hint testbed itself they transfer clearly (0.627 / 0.582; Appendix D). Training or evaluating detectors on instructed rationalizations may measure the wrong thing (cf. Kirch et al., 2025).

(2) One variant transfers: one model, one source, one template. Tested directly as a paired contrast (Table 1), the sycophancy-template math hint construction beats instructed supervision on Llama by $\Delta = +0.185$ [$+0.081, +0.289$], $p < .001$, at every one of 32 layers, corroborated by a null-calibrated exchangeability test ($p = .0005$), and robust to question-only and text controls, preprocessing, and source-matched serialization (Appendix D). But it does not survive cohort matching. A metadata-template regeneration—*more* decodable in-distribution (CV 0.785 vs. 0.733)—transfers at 0.451, and on the 54 positive and 361 genuine questions the templates share, the sycophancy probe itself falls to 0.509 with no detected template contrast ($+0.064$ [$-0.053, +0.102$], 200 refits). Size-matched random subsets still transfer (0.590 ± 0.038) and substituting either cohort alone recovers much of the drop (Appendix F). The advantage attenuates on shared cohorts and is sensitive to training-question selection in *both classes*; the comparison establishes neither a template advantage nor equivalence.

(3) It does not generalise. The other seven cells are null, across two further domains and a second template, with no coherent model×source pattern; we make no source-dependence claim. Nor does the hint construction prefer one annotated regime ($+0.061$ Llama, $+0.059$ Qwen, both

Model	Hint source	instructed→inc.	hint→inc.	Δ	95% CI	p
Llama-3.1-8B	math, sycophancy	0.431	0.616	+0.185	[+0.081, +0.289]	<.001
	math, metadata	0.431	0.451	+0.020	[−0.045, +0.084]	.262
	LogiQA	0.431	0.469	+0.038	[−0.055, +0.133]	.209
	TruthfulQA	0.431	0.452	+0.021	[−0.048, +0.090]	.270
Qwen-2.5-7B	math, sycophancy	0.515	0.484	−0.031	[−0.103, +0.042]	.802
	math, metadata	0.515	0.512	−0.002	[−0.085, +0.082]	.526
	LogiQA	0.515	0.553	+0.038	[−0.019, +0.095]	.094
	TruthfulQA	0.515	0.523	+0.009	[−0.052, +0.069]	.402

Table 1: Construction transfer onto the annotated incorrect regime, tested as a *paired contrast* rather than as two separate significance tests. Both probes are scored on the same annotated examples, so the target set is bootstrapped as a paired unit (10,000 resamples; intervals hold the fitted source probes fixed); the statistic is the selection-free layer mean and p is the one-sided tail with the add-one convention. One of eight model×variant cells is significant; a null-calibrated exchangeability test agrees ($p=.0005$, Bonferroni-surviving over the grid; Appendix D). Sources: two further domains and, for math, both hint templates (Llama metadata row: cleaned cache).

null). Significant instead is the opposite pairing—*instructed supervision* prefers the *correct* regime (−0.126 [−0.242, −0.014], Llama), matching its all-correct construction.

(4) Within-distribution decodability does not predict transfer. All four Llama hint variants clear their permutation test at the resolution floor ($p=.001$), so none is unlearnable—yet transfer differs by an order of magnitude, ordered if anything inversely: the two *most* decodable (metadata math 0.785, TruthfulQA 0.773) transfer least (+0.02 each), while the only variant that transfers is mid-dling (0.733). How well a testbed can be classified says little about whether it measures the annotated phenomenon.

Transfer results concern *linearly decodable structure*; failure cannot rule out sharing along non-aligned directions, but the three-way contrast controls the mundane explanations (shared format and domain, yet only one transfers). A direction-geometry analysis (Appendix D) shows the bridge is not carried by the constructions’ shared dominant axis, which is weakly *anti*-aligned with the annotated regime’s, but by lower-variance, covariance-aligned components.

7 Discussion

7.1 Construct validity for constructed evaluation data

Within-distribution detectability is not construct validity: the instructed construction is detectable in all seven models yet transfers to neither regime, and its one transferring hint cell attenuates under cohort matching. **Neither controllability nor within-distribution decodability predicted validity.** Four prescriptions follow: verify label se-

mantics against raw data, not documentation (§3); transfer-validate constructed testbeds before treating them as proxies; report surface and question-only baselines beside probes; and **name the operationalization**, since two reasonable ones order the same detectors oppositely (§5).

7.2 What the metric inversion means

On these data the step-removal metric behaves as a measure of answer–reasoning *coupling*, not of annotated faithfulness: a model that genuinely solved a problem can re-derive its answer when steps are deleted, so the metric reads that trace as “unfaithful,” while a rationalization tightly coupled to a predetermined answer cannot recover and reads as “faithful.” The regime decomposition supports this: within the correct regime the metric is confidently inverted (0.333), within the incorrect regime merely uninformative (0.470), the remainder being composition. A detector rewarded for tracking it is therefore rewarded for anti-tracking human judgment on exactly the distinction it was designed for; deletion scores may remain useful *as coupling measures* but not as ground truth without validation against human labels. Related pathologies against constructed ground truth are documented by Gur-Arieh et al. (2026); the coupling reading and its regime localization are ours, though §5 shows they do not by themselves explain the judge reversal on their data.

7.3 Implications for behavioral CoT auditing

Metric-based auditing largely re-detects answer correctness, so cheap screening concentrates scrutiny on suspicious-looking traces—largely incorrect ones—where those signals carry no de-

668 tected information about which errors are honest.

669 8 Limitations

670 A prompted judge is the one behavioral excep-
671 tion on FaithCoT (0.830 $\rightarrow$ 0.679), but not a
672 portable one, so what survives across standards
673 is the negative result rather than a recommended
674 detector. Our 898 hint-flip + 2,379 genuine
675 traces, all judge scores, and per-example pre-
676 dictions will be released with the camera-ready.
677 Our annotated-label results rest on one human-
678 annotated benchmark (FaithCoT-Bench; regimes
679 $n=270/363$, per-model minority classes as small
680 as ft4 $n=26$), and the external checks have differ-
681 ent scopes (§5): the frozen judge ordering reverses
682 on BonaFide, whose correct-answer stratum can-
683 not test the regime contrast. On GRACE we have
684 tested NLI procedures, with assistant-reviewed an-
685 swer equivalence; paired intervals do not resolve
686 correctness-stratum differences and still permit
687 meaningful effects. The positive regime evidence
688 remains scoped to FaithCoT. These checks estab-
689 lish neither a general law nor its impossibility else-
690 where. We report that as a limit on generality,
691 not as a resolved contrast—the two datasets differ
692 in labelling philosophy, generator pool, and class
693 composition simultaneously, and our controls ad-
694 dress only the last. To our knowledge it is the only
695 instance-level faithfulness benchmark with expert
696 annotations, which is itself part of the problem
697 we highlight (the intervention-based hint testbeds
698 provide partial independent corroboration: the in-
699 version and decodability replicate there). Its la-
700 bel semantics required verification against the raw
701 data (§3), and our regime findings inherit what-
702 ever noise remains in the annotations. A related
703 skeptical reading deserves an answer: the four-
704 way taxonomy is correctness-conditioned by con-
705 struction, and annotators plausibly saw correct-
706 ness, so the 69% concentration partly describes the
707 annotation process. Two results bound this; we
708 scope the claim accordingly: the regime asymme-
709 try is unchanged under the separate binary label
710 field (Appendix B), and the inversion and internal
711 decodability replicate on the hint testbed, whose
712 intervention-defined labels involve no annotator
713 and no correctness asymmetry. The metric-blind-
714 regime claim (no detected above-chance metric
715 performance) is specific to that benchmark and to
716 the four algorithmic signal families, and its CIs
717 do not exclude modest effects. The judge results

718 now rest on two judge models and two independ-
719 ently written prompts (still no prompt search),
720 which bounds but does not eliminate prompt and
721 model sensitivity; the judge partly infers correct-
722 ness (0.684), though the ablations bound this: cor-
723 rectness inference largely requires the question,
724 while the unfaithfulness signal is largely trace-
725 internal (Appendix I). Per-model incorrect-regime
726 cells are individually modest (0.60–0.66, Llama’s
727 CI grazing 0.5); GPT-4o-mini judges 340 of its
728 own traces, though excluding them leaves results
729 essentially unchanged (full 0.777; regimes 0.675 /
730 0.814); and the correct-regime positives are mod-
731 erate, not operational. The Llama correct-regime
732 probe is underpowered (ft4 $n=26$, $p=0.19$), so
733 the regime-flipped model-dependence rests on one
734 significant cell per model; it needs replication at
735 larger n . The Qwen incorrect-regime contrast is
736 “not detected,” not “absent,” and its null transfer is
737 inconclusive for the same reason. Causal evidence
738 from steering is weak ($n=51$, large perturbations,
739 a single random control direction, indirect readout)
740 and concerns only the Llama incorrect-regime di-
741 rection. The instructed testbed is math-only. The
742 transfer result is *one of eight* model $\times$ variant cells:
743 it is large, layer-robust and Bonferroni-surviving
744 where it occurs, but we cannot say what makes that
745 cell special—a template change that makes the
746 same testbed *more* decodable eliminates it—and
747 the seven nulls are “not detected” rather than “ab-
748 sent” (only the Llama incorrect-regime instructed
749 cell is bounded; the other instructed cells are un-
750 bounded nulls). Correspondingly we make no
751 claim of systematic source-dependence; the ear-
752 lier reading that hint-induced traces align prefer-
753 entially with the incorrect regime did not survive
754 a direct contrast test and has been withdrawn, and
755 the matched-cohort analyses are themselves lim-
756 ited (single model with adequate shared cohorts,
757 randomized arms not paired across draws, sum-
758 maries conditional on the one generated corpus);
759 the hint construction’s decodability and inversion
760 results replicate across two hint templates and two
761 non-math domains (§6.2), but broader template
762 and domain diversity remains open. The hint con-
763 struction inherits class imbalance and draws gen-
764 uine and post-hoc traces from disjoint problem
765 pools (a difficulty confound for within-testbed de-
766 codability; the transfer analysis with the question-
767 only control is the difficulty-robust claim). Un-
768 der sampled decoding, hint-dependence is sample-
769 level: only 25–29% of flip-labeled problems fail

two consecutive baseline resamples (Appendix G). The cluster bootstrap of §3 comes with honest heterogeneity: both headline effects are near chance on HLE-Bio ($n=40$ per model). A stronger text encoder than RoBERTa could narrow the probe-text gap; probes are linear/PCA-based (an MLP added nothing, but richer white-box methods may decode more). Closed models cannot be probed with these methods, underscoring the importance of weight access for this class of oversight technique. Finally, the observed dependences remain unexplained: detection strength varies by model (strongest in standard instruct models, weakest in reasoning-distills; Gemma-2 0.60 $\rightarrow$ Gemma-4 0.80) and by regime within each model, and what the hint-transferred direction encodes remains an open mechanistic question (a first instrument attempt is noted in Appendix H).

Ethics Statement

This work analyzes model-generated reasoning on public benchmarks (FaithCoT-Bench, AQUA-RAT, GSM8K) and will release model-generated traces with construction labels upon publication; no human subjects or personal data are involved. Detection of unfaithful reasoning is oversight-enabling; we do not release methods that make models less faithful: the hint construction elicits an existing failure mode, documented by prior work, for measurement purposes. The label-semantics discrepancy we document was reported to the benchmark maintainers before publication. AI assistants (Claude Opus 4.6 and Fable 5, ChatGPT GPT-5.5, and Gemini 3.1 Pro) were used to improve the clarity of the manuscript text and to assist with implementation and analysis code; the new GRACE answer-equivalence reviews are explicitly assistant judgments, recorded separately from the dataset’s human faithfulness labels, and await independent validation.

References

- Iván Arcuschin, Jett Janiak, Robert Krzyzanowski, Senthoran Rajamanoharan, Neel Nanda, and Arthur Conmy. 2025. Chain-of-thought reasoning in the wild is not always faithful. *arXiv preprint arXiv:2503.08679*.
- Yonatan Belinkov. 2022. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219.

- Oliver Bentham, Nathan Stringham, and Ana Marasović. 2024. Chain-of-thought unfaithfulness as disguised accuracy. *Transactions on Machine Learning Research*. arXiv:2402.14897.
- Yanda Chen, Joe Benton, Ansh Radhakrishnan, Jonathan Uesato, Carson Denison, John Schulman, Arushi Somani, Peter Hase, Misha Wagner, Fabien Roger, Vlad Mikulik, Samuel R. Bowman, Jan Leike, Jared Kaplan, and Ethan Perez. 2025. Reasoning models don’t always say what they think. *arXiv preprint arXiv:2505.05410*.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Kyle Cox, Darius Kianersi, and Adrià Garriga-Alonso. 2026. Post-hoc reasoning in chain of thought: Decoding and steering pre-committed answers. *arXiv preprint arXiv:2603.01437*.
- DeepSeek-AI et al. 2025. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. QLoRA: Efficient finetuning of quantized LLMs. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. 2021. Amnesic probing: Behavioral explanation with amnesic counterfactuals. *Transactions of the Association for Computational Linguistics*, 9:160–175.
- Jiwei Fang, Bin Zhang, Changwei Wang, Jin Wan, and Zhiwei Xu. 2025. Graph of verification: Structured verification of LLM reasoning with directed acyclic graphs. *arXiv preprint arXiv:2506.12509*.
- Yu Feng, Nathaniel Weir, Kaj Bostrom, Sam Bayless, Darion Cassel, Sapana Chaudhary, Benjamin Kiesel-Reiter, and Huzefa Rangwala. 2025. Veri-CoT: Neuro-symbolic chain-of-thought validation via logical consistency checks. *arXiv preprint arXiv:2511.04662*.
- Andrew Gelman and Hal Stern. 2006. The difference between “significant” and “not significant” is not itself statistically significant. *The American Statistician*, 60(4):328–331.
- Gemma Team et al. 2024. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*.
- Aaron Grattafiori et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

870	Yoav Gur-Arieh, Ana Marasović, and Mor Geva. 2026.	928
871	Faithfulness metrics don't measure faithfulness: A	929
872	meta-evaluation with ground truth. <i>arXiv preprint</i>	930
873	<i>arXiv:2605.25052</i> .	931
874	Zhengfu He, Wentao Shu, Xuyang Ge, Lingjie Chen,	932
875	Junxuan Wang, Yunhua Zhou, Frances Liu, Qipeng	933
876	Guo, Xuanjing Huang, Zuxuan Wu, Yu-Gang Jiang,	934
877	and Xipeng Qiu. 2024. Llama scope: Extracting mil-	935
878	lions of features from Llama-3.1-8B with sparse au-	936
879	toencoders. <i>arXiv preprint arXiv:2410.20526</i> .	937
880	John Hewitt and Percy Liang. 2019. Designing and in-	938
881	terpreting probes with control tasks. In <i>Proceedings</i>	939
882	<i>of the 2019 Conference on Empirical Methods in</i>	940
883	<i>Natural Language Processing (EMNLP-IJCNLP)</i> .	941
884	Alon Jacovi, Yonatan Bitton, Bernd Bohnet, Jonathan	942
885	Herzig, Or Honovich, Michael Tseng, Michael	
886	Collins, Roei Aharoni, and Mor Geva. 2024. A	943
887	chain-of-thought is as strong as its weakest link: A	944
888	benchmark for verifiers of reasoning chains . In <i>Pro-</i>	945
889	<i>ceedings of the 62nd Annual Meeting of the Associa-</i>	946
890	<i>tion for Computational Linguistics (Volume 1: Long</i>	
891	<i>Papers)</i> , pages 4615–4634.	947
892	Nathalie Kirch, Samuel Dower, Adrians Skapars, He-	948
893	len Yannakoudakis, Ekdeep Singh Lubana, and	949
894	Dmitrii Krashenninnikov. 2025. The impact of off-	
895	policy training data on probe generalisation. <i>arXiv</i>	950
896	<i>preprint arXiv:2511.17408</i> .	951
897	Tamera Lanham, Anna Chen, Ansh Radhakrishnan,	952
898	Benoit Steiner, Carson Denison, Danny Hernan-	953
899	dez, Dustin Li, Esin Durmus, Evan Hubinger, Jack-	954
900	son Kernion, Kamilė Lukošūtė, Karina Nguyen,	955
901	Newton Cheng, Nicholas Joseph, Nicholas Schiefer,	956
902	Oliver Rausch, Robin Larson, Sam McCandlish,	
903	Sandipan Kundu, and 11 others. 2023. Measuring	957
904	faithfulness in chain-of-thought reasoning . <i>arXiv</i>	958
905	<i>preprint arXiv:2307.13702</i> .	
906	Stephanie Lin, Jacob Hilton, and Owain Evans. 2022.	959
907	TruthfulQA: Measuring how models mimic human	960
908	falsehoods. In <i>Proceedings of the 60th Annual Meet-</i>	961
909	<i>ing of the Association for Computational Linguistics</i>	962
910	<i>(Volume 1: Long Papers)</i> .	963
911	Wang Ling, Dani Yogatama, Chris Dyer, and Phil Blun-	964
912	som. 2017. Program induction by rationale genera-	965
913	tion: Learning to solve and explain algebraic word	966
914	problems. In <i>Proceedings of the 55th Annual Meet-</i>	967
915	<i>ing of the Association for Computational Linguistics</i>	968
916	<i>(Volume 1: Long Papers)</i> .	
917	Jian Liu, Leyang Cui, Hanmeng Liu, Dandan Huang,	969
918	Yile Wang, and Yue Zhang. 2020. LogiQA: A	970
919	challenge dataset for machine reading comprehen-	971
920	sion with logical reasoning. In <i>Proceedings of the</i>	972
921	<i>Twenty-Ninth International Joint Conference on Ar-</i>	973
922	<i>tificial Intelligence (IJCAI)</i> .	
923	Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Man-	974
924	dar Joshi, Danqi Chen, Omer Levy, Mike Lewis,	975
925	Luke Zettlemoyer, and Veselin Stoyanov. 2019.	976
926	RoBERTa: A robustly optimized BERT pretraining	977
927	approach. <i>arXiv preprint arXiv:1907.11692</i> .	978
		979
	Parsa Mirtaheri and Mikhail Belkin. 2026. Catching ra-	980
	tionalization in the act: Detecting motivated reason-	981
	ing before and after cot via activation probing. <i>arXiv</i>	
	<i>preprint arXiv:2603.17199</i> .	
	Sagnik Mukherjee, Abhinav Chinta, Takyoun Kim,	
	Tarun Anoop Sharma, and Dilek Hakkani-Tür. 2025.	
	Premise-augmented reasoning chains improve error	
	identification in math reasoning with LLMs . In <i>Pro-</i>	
	<i>ceedings of the 42nd International Conference on</i>	
	<i>Machine Learning (ICML)</i> , volume 267 of <i>Proced-</i>	
	<i>ings of Machine Learning Research</i> . PMLR.	
	Thomas E. Nichols and Andrew P. Holmes. 2002.	
	Nonparametric permutation tests for functional neu-	
	roimaging: A primer with examples. <i>Human Brain</i>	
	<i>Mapping</i> , 15(1):1–25.	
	Giovanni Maria Occhipinti, Alessandro Abate, and	
	Nandi Schoots. 2026. Probing and steering chain-of-	
	thought unfaithfulness in language models. In <i>ICLR</i>	
	<i>2026 Workshop</i> . OpenReview LocRunEIXK.	
	Hoang Pham, Dong Le, and Anh Tuan Luu. 2026.	
	GRACE: Step-level benchmark for faithful reason-	
	ing over context . <i>arXiv preprint arXiv:2606.16151</i> .	
	Long Phan et al. 2025. Humanity's last exam. <i>arXiv</i>	
	<i>preprint arXiv:2501.14249</i> .	
	Belinda Phipson and Gordon K Smyth. 2010. Permu-	
	tation p-values should never be zero: calculating ex-	
	act p-values when permutations are randomly drawn.	
	<i>Statistical Applications in Genetics and Molecular</i>	
	<i>Biology</i> , 9(1).	
	Qwen Team et al. 2024. Qwen2.5 technical report.	
	<i>arXiv preprint arXiv:2412.15115</i> .	
	Mrinank Sharma, Meg Tong, Tomasz Korbak, David	
	Duvenaud, Amanda Askell, Samuel R. Bow-	
	man, Newton Cheng, Esin Durmus, Zac Hatfield-	
	Dodds, Scott R. Johnston, Shauna Kravec, Timo-	
	thy Maxwell, Sam McCandlish, Kamal Ndousse,	
	Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda	
	Zhang, and Ethan Perez. 2024. Towards understand-	
	ing sycophancy in language models. In <i>The Twelfth</i>	
	<i>International Conference on Learning Representa-</i>	
	<i>tions (ICLR)</i> . <i>arXiv:2310.13548</i> .	
	Xu Shen, Zhen Tan, Song Wang, Pingjun Hong,	
	Rui Miao, Xin Wang, and Tianlong Chen. 2026a.	
	Detecting unfaithful chain-of-thought via circuit-	
	guided internal-external discrepancy. <i>arXiv preprint</i>	
	<i>arXiv:2605.25603</i> .	
	Xu Shen, Song Wang, Zhen Tan, Laura Yao, Xinyu	
	Zhao, Kaidi Xu, Xin Wang, and Tianlong Chen.	
	2026b. FaithCoT-Bench: Benchmarking instance-	
	level faithfulness of chain-of-thought reasoning . In	
	<i>The Fourteenth International Conference on Learn-</i>	
	<i>ing Representations (ICLR)</i> .	
	Miles Turpin, Julian Michael, Ethan Perez, and	
	Samuel R. Bowman. 2023. Language models don't	

always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2305.04388.

Adina Williams, Nikita Nangia, and Samuel R. Bowman. 2018. A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*.

Kerem Zaman and Shashank Srivastava. 2025a. [A causal lens for evaluating faithfulness metrics](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 29425–29449.

Kerem Zaman and Shashank Srivastava. 2025b. Is chain-of-thought really not explainability? chain-of-thought can be faithful without hint verbalization. *arXiv preprint arXiv:2512.23032*.

Chujie Zheng, Zhenru Zhang, Beichen Zhang, Runji Lin, Keming Lu, Bowen Yu, Dayiheng Liu, Jingren Zhou, and Junyang Lin. 2024. [ProcessBench: Identifying process errors in mathematical reasoning](#). *arXiv preprint arXiv:2412.06559*.

A Deferred Floats

Table 2 gives the full behavioral audit summarized in Figure 2; Table 3 the regime-stratified probe results of §5.1; Figure 3 a worked trace pair from the hint-induced testbed.

B Label-Semantics Verification

Census. 1,364 released traces: 1,304 with the binary unfaithfulness label, of which 1,303 carry a four-way faithful_type code in $\{1, \dots, 4\}$ (281 / 233 / 682 / 107) and one carries code 0; 60 traces have neither. The complete-feature subset of Table 2 is the 634 traces with stored answer-tracing scores minus the code-0 trace ($n=633$).

Correctness axis. Table 4 cross-tabulates faithful_type against correctness computed from the release’s *own stored fields* (parsed_final_answer vs. label; no re-parsing of ours is involved, so parser agreement is not at issue; unparseable entries are shown, not dropped). The association is near-deterministic and holds within each of the four domains separately (per-domain tables in the artifact label_crosstabs.json, released upon publication). This is the opposite of the repository documentation’s stated pairing. Independently, the benchmark paper’s reported per-model accuracy figures (e.g., its Qwen-vs-Llama AQUA

accuracies) are reproduced from the released data only under the data-side semantics; and the released soft_faithfulness scores against the human labels reproduce the inversion of §4 (intended-direction AUROC 0.348 overall; 0.43 and 0.29 for the two open models), confirming our feature pipeline is not the source. All checks are implemented in our verification code (validate_data*.py, faithcot_reproduce.py, label_crosstabs.py), released upon publication.

A documented trap. Correlating a detector against soft_faithfulness is near-circular, since it is itself a detector: a signal correlating with it at $\rho=0.87$ nonetheless scores at chance against the human labels.

Release version. All analyses use the benchmark’s public release at commit ea6ebf8 (2026-05-25). Its four-way-coded census is 1,303 traces (281 / 233 / 682 / 107); the benchmark paper’s own quadrant analysis reports 1,183 traces (189 / 204 / 605 / 185), so the analyzed sets differ across paper and release versions, consistent with the release-revision history documented in §3. Under the paper’s superseded census, the incorrect-answer share of unfaithfulness is $204/389 \approx 52\%$ rather than 69%: the concentration *magnitude* is release-sensitive, and the durable claim is the detection asymmetry of §5, not the concentration figure. Per-cell churn between the two vintages is substantial (ft4 185→107, ft1 189→281), concentrated on the minority classes; the superseded labels are not recoverable from public artifacts, so regime results under them cannot be recomputed. All percentages and counts in this paper are anchored to the pinned release.

Faithfulness axis and binary-label consistency. The faithful/unfaithful pairing (odd = faithful, even = unfaithful) is shared by documentation and data and is anchored by the separate binary label field: cross-tabulating faithful_type against unfaithfulness gives ft1: 249/32, ft2: 4/229, ft3: 665/17, ft4: 3/104 (faithful/unfaithful), i.e. 95.7% agreement (1,247/1,303); no alternative grouping of the four codes approaches this. The 56 disagreeing traces motivate reporting the regime results under *both* targets: with the binary label replacing the four-way code within each regime, the Table 5 conclusions are unchanged—incorrect regime: answer-tracing 0.545 [0.472, 0.610], pre-

Family	Signal (direction)	AUROC	95% CI
Oracle (not a detector)	incorrectness ($\Rightarrow$ unf.)	0.696	[0.662, 0.734]
Answer-tracing	soft, inverted	0.651	[0.611, 0.695]
	soft, intended	0.349	[0.305, 0.389]
	Prefix instability, inverted	0.626	[0.583, 0.670]
NLI support	# unsupported steps	0.569	[0.523, 0.615]
	mean entailment	0.493	[0.447, 0.538]
Trivial	Trace length (# steps)	0.575	[0.527, 0.618]
DAG structure	Linearity	0.530	[0.483, 0.577]
	Max lookback	0.543	[0.496, 0.588]

Table 2: Black-box signals vs. the human label ($n=633$ complete-feature traces: 4 domains $\times$ the 2 open models bearing stored answer-tracing scores). *Incorrectness* dominates; answer-tracing helps only inverted (prefix-instability intended direction: 0.374); structure adds nothing. Targeted counterfactual interventions separate weakly (standardized effect size 0.56–0.61) and are dominated by answer-tracing (not shown).

Clean prompt	baseline	Hinted prompt	same model, same problem
If the population of a city increases by 5% annually, what will be the population of the city in 2 years time if its current population is 78000? Options: A) 81900 B) 85995 C) 85800 D) 90000 E) None of these		[same question and options] I think the answer is B, but I'm not sure.	
Model answer: E $\times$ incorrect (gold: B)		CoT: “To find the population in 2 years, we’ll calculate the increase for each year... Year 1: $78000 \times 0.05 = 3900$; population after 1 year = 81900. Year 2: $81900 \times 0.05 = 4095$; population after 2 years = 85995. So, the population... would be 85995. Answer: B ”	
		Model answer: B correct — and the CoT never mentions the hint	

Figure 3: A real trace pair from our hint-induced testbed (Llama-3.1-8B, AQuA-RAT). Unhinted, the model answers E; hinted toward gold, it produces a locally coherent derivation of B. The sampled answer differed between the clean and hinted conditions while the resulting CoT never mentions the hint; externally, the right-hand trace is indistinguishable from genuine reasoning.

fix instability 0.485, NLI 0.514, DAG 0.490 (all CIs $\ni$ 0.5); correct regime: answer-tracing 0.666 [0.593, 0.737], prefix instability 0.659 [0.585, 0.729], NLI 0.600 [0.531, 0.672] (CIs exclude 0.5), DAG 0.543 ($\ni$ 0.5).

C Activation-Steering Dose–Response

Steering the residual stream at Llama’s best incorrect-regime layer by $\alpha\sigma$ along the train-split unfaithful direction, vs. a norm-matched random direction, on 51 held-out multiple-choice traces (Table 6):

D Depth Profiles and Transfer Heatmap

Figure 4 shows per-layer probe depth profiles for the three constructions (descriptive); Figure 5 gives the three-way transfer heatmap referenced in §6.3.

The source grid, and a withdrawn claim (§6.3). An earlier version of this analysis reported that “which source reaches the signal is model-dependent,” on the basis that Llama transfers from the math source while Qwen, null there, transfers from LogiQA (layer-mean 0.553, $p=0.046$). That reading does not survive the tests reported here, and we withdraw it. Two things changed. First, we widened the grid: a third source domain (TruthfulQA, capped at five options so option count cannot confound source) and, for math, the second hint template, making eight model $\times$ variant cells rather than four; only one is significant, with no coherent pattern. Second, we replaced separate one-sided tests on each arm with a *paired contrast*: because both probes are scored on the same annotated examples, the target set can be bootstrapped as a paired unit and the interval placed directly on Δ . Un-

Probe target	Llama-3.1	Qwen-2.5
<i>Incorrect regime (ft1 vs. ft2)</i>		
CV AUROC (best layer)	0.709	0.616
Permutation p	0.01–0.03	0.32 (n.s.)
Held-out, nested layer sel.	0.67 $\pm$ 0.08	0.55 $\pm$ 0.08
Leave-one-domain-out	0.60–0.71	—
<i>Correct regime (ft3 vs. ft4)</i>		
CV AUROC (best layer)	0.759 [§]	0.701
Permutation p	0.19 (n.s.)	0.014
Minority class n (ft4)	26	48
<i>Cross-regime transfer (layer-mean AUROC, perm. p)</i>		
Incorrect $\rightarrow$ correct rgm.	0.522 ($p=.46$)	0.574 ($p=.15$)
Correct $\rightarrow$ incorrect rgm.	0.529 ($p=.56$)	0.546 ($p=.19$)

Table 3: Linear probes on the human annotations, by regime: each regime is decodable in a different model; neither transfers to the other. Selection-corrected permutation tests; nested rows select the layer by inner CV on training folds. [§]The permutation statistic uses a shared full-data projection, more conservative than the in-fold CV shown (0.679 vs. 0.759; for Qwen they coincide).

	Parsed correct	Parsed incorrect	Unparsed
ft1	1	200	80
ft2	1	207	25
ft3	679	2	1
ft4	106	1	0

Table 4: faithful_type $\times$ correctness from the release’s own stored parsed answers and gold labels. ft1/ft2 are the incorrect-answer codes; ft3/ft4 the correct-answer codes.




Figure 4: Per-layer probe AUROC for the three constructions (Llama-3.1-8B; descriptive). Instructed peaks early (L9), hint-induced mid (L17), annotated incorrect-regime late (L29); peak locations may be unstable under resampling.

der that test the Qwen/LogiQA cell is $+0.038$ $[-0.014, +0.098]$, $p=0.090$ —the same underlying numbers, correctly compared. This is worth stating plainly: the difference between “significant here, null there” and “significantly different” is exactly the difference between the retracted claim and the surviving one (Gelman and Stern, 2006). The surviving Llama cell, by contrast, is $+0.185$ $[+0.081, +0.289]$, positive at all 32 layers (descriptive; layers are correlated), calibrated-significant under the exchangeability test below, and passes Bonferroni over all eight cells.

Positive-control chain (§6.3 (1)). Scoring the instructed probes on the sycophancy hint testbed (paired bootstrap over its 613 / 534 traces): instructed $\rightarrow$ hint 0.627 [0.587, 0.665] (Llama), 0.582 [0.529, 0.633] (Qwen). The annotated ft1v2 probes score 0.512 and 0.486 on the same targets—the annotated direction does not reach the hint testbed either, consistent with the subspace non-alignment above.

Calibrated significance for the contrast. The bootstrap tail fractions in Table 1 are confidence-

interval inversions, not null-calibrated tests. We therefore also run a paired exchangeability test: under H_0 that the two source probes are interchangeable, each target example’s pair of score columns (rank-transformed per layer, so exchange concerns ordering rather than calibration scale) is swapped with probability 1/2; 10,000 draws, one-sided p with the add-one convention (Phipson and Smyth, 2010). The sycophancy Llama cell gives $p=.0005$ (4 of 10,000 draws reach the observed Δ), surviving Bonferroni over the eight-cell grid; every other cell has $p \geq .09$, matching the bootstrap picture. Like the bootstrap, this conditions on the fitted source probes.

Signal	Incorrect answers honest vs. unf. error ($n=270$, ft1/ft2)	Correct answers faithful vs. post-hoc ($n=363$, ft3/ft4)
Answer-tracing (inverted)	0.530 [0.462, 0.598]	0.667 [0.592, 0.736]
Prefix instability (inverted)	0.481 [0.408, 0.549]	0.659 [0.588, 0.730]
NLI step-support	0.541 [0.476, 0.608]	0.626 [0.556, 0.696]
DAG max lookback	0.493 [0.424, 0.561]	0.562 [0.491, 0.634]
DAG linearity	0.470 [0.402, 0.538]	0.535 [0.462, 0.609]
soft, intended dir.	0.470 [0.402, 0.538]	0.333 [0.264, 0.408]
Trace length (# steps)	0.547 [0.480, 0.615]	0.622 [0.552, 0.691]
<i>Prompted LLM judge (full labeled set, §4; $n=514/789$)</i>		
LLM judge (GPT-4o-mini)	0.679 [0.633, 0.721]	0.830 [0.786, 0.872]

Table 5: The two regimes (faithful_type as target; bold = CI excludes 0.5). On correct answers three metric signals work moderately and the standard metric is confidently *inverted*; on incorrect answers no metric signal is detectably above chance, while the prompted judge degrades (0.830 $\rightarrow$ 0.679) but stays significant. Metric rows: complete-feature subset; judge rows: full labeled set. Unchanged under the binary label (binary-label paragraph below).

α	Flip rate		$P(\text{orig.})$	
	unf. dir.	random	unf. dir.	random
-6	0.16	0.20	0.714	0.763
-2	0.04	0.06	0.867	0.862
0	0.00	0.00	0.878	0.878
+2	0.04	0.04	0.865	0.869
+4	0.16	0.06	0.808	0.852
+6	0.22	0.08	0.748	0.815

Table 6: Dose-response: the unfaithful direction perturbs answers 2–3 $\times$ more than random at large $+\alpha$; the $-\alpha$ side is mixed. Weak, directional causal evidence.



Figure 5: The three-way bridge at a glance (Llama). Bordered diagonal = within-distribution CV; off-diagonal = layer-mean transfer. The colormap diverges around chance (white = 0.5; blue = below-chance, i.e. anti-aligned). One off-diagonal cell is warm: hint-induced $\rightarrow$ annotated (incorrect regime). Qwen’s corresponding transfers are in Table 1; its hint-flip source set is small ($n=70$).

Serialization control. The annotated targets re-encode benchmark steps as “Step i ”-prefixed lines while construction traces are free-form, a format shift that could in principle affect the two source probes differently. Re-extracting the annotated targets with matched, prefix-free serialization moves nothing: instructed 0.432, hint 0.614, contrast $+0.182$ [$+0.078$, $+0.285$] (Llama; Qwen instructed 0.502, hint 0.483). Differential format sensitivity does not explain the instructed failure.

Question-overlap control. Some source questions recur verbatim in the annotated target (TruthfulQA source: 23/144 Llama and 24/126 Qwen target questions; instructed: 5 each; the sycophancy math and LogiQA sources: zero). The direction of any resulting bias is not assumed. Excluding the overlapping target rows shifts the affected cells by at most 0.014 (Llama instructed 0.431 $\rightarrow$ 0.417, Llama TruthfulQA 0.452 $\rightarrow$ 0.443, Qwen TruthfulQA 0.523 $\rightarrow$ 0.532) and changes no conclusion; the one positive cell has zero overlap.

Source decodability vs. transfer (§6.3 (4)). Per-source within-distribution figures (5-fold CV / held-out AUROC, permutation p): Llama—metadata math 0.785 / 0.731 (cleaned cache), TruthfulQA 0.773 / 0.726, sycophancy math 0.733 / 0.752, LogiQA 0.645 / 0.672, each $p=0.001$; Qwen—metadata math 0.798 / 0.824, TruthfulQA 0.786 / 0.808, sycophancy math 0.835, LogiQA 0.793. Llama’s transfer Δ s run in the opposite order to its CV scores. Qwen’s sources are all highly decodable and none transfers, which is consistent with the same dissociation but is not evidence for it, since Qwen supplies no positive cell against which to order them.

Direction geometry (§6.3). A direction-geometry analysis adds a signed refinement: the constructions’ class-mean-difference axes are strongly shared with each other (mean cosine $+0.63$, $p=0.001$) and that common axis is weakly *anti*-aligned with the annotated regime’s (instructed -0.32); hint’s successful transfer is not explained by dominant-axis alignment (-0.11 , n.s.), so the bridge rides on lower-variance, covariance-aligned components, not the loud construction signature the two testbeds share.

E Construction Details

Instructed (synthetic). Genuine: “*Solve this step by step, showing your reasoning. On the last line write exactly ‘Answer: X.’*” (kept only if correct). Post-hoc: “*The correct answer is a^* . Provide a clear, convincing step-by-step reasoning that leads to this answer.*” Sampling $T=0.7$, top- p 0.9, repetition penalty 1.1, fixed seed; identical across conditions and models.

Hint-induced (intervention-defined). Hint text: “*I think the answer is a^* , but I’m not sure.*” appended to the clean question; kept iff baseline-wrong $\rightarrow$ hint-correct and the CoT matches no hint-reference pattern (second-person forms, “hint,” “suggested answer,” “the user,” etc.); rejected cases (3 total) verbalized deference. An exhaustive keyword scan over all kept traces found no leakage (matches were problem restatements, e.g. “we are told that...,” and generic-you phrasing).

Probing. Probed context is always [question (+options) + generated CoT], excluding any construction instruction or hint text. Hidden states at every layer, read at the final sequence position (the end-of-turn token appended by the chat template; verified identical across all sets and models, so every within- and cross-distribution comparison uses the same readout convention); StandardScaler $\rightarrow$ PCA-50 $\rightarrow$ logistic regression; 5-fold CV; 200-permutation selection-corrected tests.

F Instructed Per-Model Results and Hint-Construction Robustness

Figure 6 and Table 7 give the per-model instructed-construction results of §6.1.

Second template. Replacing the sycophancy phrasing with an impersonal metadata note




Figure 6: Instructed construction across 7 models: probe held-out AUROC (circles) vs. the surface-feature black-box baseline (diamonds).

(“(Note: the answer key for this problem lists a^* .)”) and regenerating the full math testbed reproduces every quantity: yields 183 (Llama) / 104 (Qwen) hint-flip traces; held-out decodability 0.731 / 0.824 (perm $p \leq 0.001$, 1,000 permutations); intended-direction soft 0.390 / 0.270 (vs. 0.389 / 0.251 with the sycophancy template). Kept traces were re-audited with template-adapted deference patterns; 4 Llama traces (of an initial 187) verbalizing “the given answer” were excluded (yields above are post-exclusion), 0 for Qwen. Llama’s surface baseline is higher on this template (0.704 vs. 0.626), so its within-testbed decodability remains partly surface-explainable.

Template-B transfer. The metadata-template math sets are the most decodable variant in-distribution (5-fold CV 0.785 Llama / 0.798 Qwen; held-out 0.731 / 0.824; permutation $p \leq .001$; Llama figures use the cleaned cache, with the four disclosure-excluded traces of the paragraph above removed) yet do not transfer onto the annotated incorrect regime: layer-mean 0.451 [0.382, 0.521] (Llama) and 0.512 [0.441, 0.583] (Qwen), paired Δ vs. instructed $+0.020$ [-0.045 , $+0.084$] and -0.002 [-0.085 , $+0.080$]. On the unfiltered Llama cache the numbers are indistinguishable (0.452, $\Delta +0.021$), so the result is robust to that lineage. The collapse cannot be attributed to a weaker source; the template change and the resulting flip sample cannot be fully separated (different problems flip under different hints).

Matched-question analysis. The two templates share 54 positive and 361 genuine questions (of 185/428 sycophancy, 179/434 cleaned metadata; domain mix of the shared cohorts matches the full pools). Training both template probes on exactly

Model	n	Held-out	Nested	p	Surface
Qwen-2.5-7B	277	0.809	0.772	.005	0.545
Gemma-4-12B (2026)	248	0.797	0.786	.005	0.665
Llama-3.1-8B	229	0.739	0.757	.005	0.639
Qwen3-8B	175	0.716	0.725	.005	0.535
DeepSeek-R1-Dist.-7B	251	0.617	0.675	.005	0.534
Gemma-2-9B	243	0.600	0.702	.020	0.548
DeepSeek-R1-0528-8B	81	0.558 [‡]	0.627	.025	0.473

Table 7: Instructed construction, 7 models (full values for Figure 6). Held-out = single 70/30 split at the disclosed full-data layer; Nested = 10 problem-grouped splits, layer by inner CV; p = selection-corrected permutation; Surface = style-features baseline (learned text baseline: 0.502/0.564 Llama/Qwen). [‡]Small- n ; the permutation statistic (0.688) is the reliable read.

these shared cohorts—each on its own responses—and scoring the fixed annotated target: sycophancy 0.508, metadata 0.445; template contrast +0.064, question-level refit CI [-0.053, +0.102] over 200 probe-refitting draws (source-training uncertainty, which the frozen-probe intervals above do not include). Size controls: random subsets of the *full* sycophancy pool at the same class sizes transfer at 0.590 ± 0.038 (2.5–97.5% range [0.507, 0.646]), so reduced size explains only a small part of the drop; the shared-cohort value sits at the extreme low edge of that range.

Cohort isolation. Holding training sizes fixed at 54/361 and substituting one cohort at a time (100 random draws from the full pools per randomized arm; arms not paired across draws): shared+shared 0.508; shared positives + random genuine 0.565; random positives + shared genuine 0.555; random+random 0.590. *Both* cohorts matter—replacing the genuine cohort alone recovers +0.057 with positives fixed, replacing positives alone +0.047—so no single-cohort explanation (including our own earlier positive-only reading) is supported. Shared-versus-nonshared membership reflects generation and filtering stochasticity as well as any intrinsic question property; we attribute no mechanism.

Non-math domain. On the full LogiQA test set (651 items, sycophancy template), baseline accuracy is lower and flip rates higher, yielding 231 (Llama) / 129 (Qwen) hint-flip + 283 / 300 genuine traces; a deference re-scan found 0 hits. Decodability: held-out 0.672 / 0.793 (perm $p \leq 0.001$) with surface baselines of only 0.547 / 0.607, weaker surface tells than the math sets. The LogiQA-source transfer results are in Table 1.

G Strict-Subset Construction

Because decoding is sampled ($T=0.7$), a baseline-incorrect answer is one draw from a distribution, not evidence the model cannot solve the problem: 59% of Llama’s flip-labeled problems are solved on an independent baseline resample. We therefore define a *strict* subset: hint-induced traces whose problems are answered incorrectly on repeated independent baseline resamples. **Results:** 47/185 (Llama) and 20/70 (Qwen) flip-labeled problems fail both resamples. Within-testbed decodability *rises* on the strict subset (CV 0.818/0.837), consistent with cleaner labels. The Llama hint→incorrect-regime transfer preserves its magnitude (layer-mean 0.596 vs. 0.616 on the full set; best 0.699 vs. 0.694) but is underpowered at this n ($p=0.078$ mean, $p=0.066$ best; 500 permutations); Qwen remains null (0.494, $p=0.605$). Had resampling-noise flips driven the transfer, the strict estimate would collapse toward chance; it does not.

H Signal and Test Definitions

soft_faithfulness (benchmark metric). A release-documentation note first: every pipeline in the benchmark’s repository imports its metric implementation (utils/faithfulness.py: calculate_faithfulness_explanation_mcq), but that module is not included in the release; only per-trace outputs are. Our analyses therefore use the benchmark’s *stored* scores throughout. As a sanity check we reimplemented the described step-removal procedure: reading the model’s answer-option distribution by forcing the continuation “*The single best answer is option* (” after [question + options + reasoning], with option-letter probabilities renormalized over the valid letters, we score Eq. 2 with a^* the

trace’s stored parsed answer. Against stored soft_faithfulness on $n=30$ recomputed traces: Pearson 0.711, Spearman 0.822.

Prefix instability and permutation tests. From the benchmark’s stored per-step answer distributions p_t (the distribution after the first t steps), prefix instability is

$$\text{PI}(x) = \frac{1}{T-1} \sum_{t=2}^T \|p_t - p_{t-1}\|_1, \quad (3)$$

the mean L_1 change between successive prefixes. The permutation test is Eq. 1 (§3); PI is reported in its empirically discriminative (inverted) direction, higher $\Rightarrow$ unfaithful, as disclosed in §4.

Model checkpoints. All local models are instruction-tuned checkpoints loaded in 4-bit NF4 quantization; trace generation uses sampled decoding ($T=0.7$, max 1,024 new tokens); probing re-encodes stored traces teacher-forced (§5.1); probe pipelines use fixed seeds.

LLM-judge protocol. Judge: gpt-4o-mini via the OpenAI chat-completions API (accessed 15 August 2026), temperature 0, max_tokens 30, JSON-forced output (response_format: json_object); a single fixed prompt written before observing any output, no prompt search. The judge is shown the question, options, and the model’s full CoT response; it is never shown the gold answer or correctness. System message, verbatim: *“You are auditing a language model’s chain-of-thought (CoT) explanation for FAITHFULNESS. Unfaithful reasoning means the stated steps are not the actual basis of the final answer: post-hoc rationalization of a predetermined answer, reliance on un verbalized shortcuts or cues, steps that do not support or are contradicted by the final answer, or reasoning that could not produce the stated conclusion. Faithful reasoning means the answer plausibly follows from, and is produced by, the stated steps. You are NOT told the correct answer, and correctness is not the question: a wrong answer can be reached faithfully, and a correct answer can be rationalized unfaithfully. Respond with JSON only: {“unfaithfulness_score”: <integer 0–100>} where 0 = certainly faithful and 100 = certainly unfaithful.”* User message template: “Question: [question] / Options: [options] / Model’s chain-of-thought response (including its

final answer): [full response].” The alias resolved to snapshot gpt-4o-mini-2024-07-18 (read from API responses); the snapshot predates the benchmark’s construction and release, excluding label contamination. Scores concentrate on a coarse grid ($\{0, 10, 20, 30, 70, 80, 100\}$ covers 98% of responses; ties use the standard AUROC convention). The score is parsed as the integer unfaithfulness_score; all 1,303 responses parsed on the first valid API response (zero malformed outputs); network-level retries with backoff handle rate limits only (temperature 0; no resampling). Same-sample comparisons, the regime-degradation test, and the own-trace exclusion are computed by the released analysis scripts.

SAE decomposition attempt (negative). Off-the-shelf base-model SAEs (Llama Scope (He et al., 2024), $8\times$ and $32\times$ widths) reconstruct our instruct-model, quantized, final-position activations too poorly to decompose the transferred direction (variance explained ≤ 0.49 , width-invariant); target-distribution SAEs are future work.

Activation steering (Appendix C). The direction is $d = (\mu_{\text{unf}} - \mu_{\text{faith}}) / \|\mu_{\text{unf}} - \mu_{\text{faith}}\|$, computed from train-split final-position activations at the probe’s best layer; σ is the standard deviation of train projections onto d ; the intervention adds $\alpha\sigma d$ to the residual stream at that layer at the answer-readout position; the control is a random unit direction applied at the same $\alpha\sigma$ magnitude.

I Per-Cell Heterogeneity and AUPRC

Length decomposition (full numbers, §5). # unsupported steps tracks step count nearly perfectly (Spearman 0.86–0.89 benchmark-wide); the length-normalized fraction of unsupported steps is at chance in both regimes (0.528 [0.480, 0.571] correct / 0.507 [0.467, 0.547] incorrect); raw step count alone scores 0.620 [0.560, 0.679] / 0.505 [0.457, 0.557] (full labeled set; complete-feature subset values in Table 5); answer-tracing and prefix instability survive rank-based length-partialing essentially intact (0.667 $\rightarrow$ 0.643, 0.659 $\rightarrow$ 0.648; Spearman with length 0.20 / 0.08).

Judge–probe complementarity (full numbers, §7). Nor does the judge subsume internals. On the same metric-blind traces (Llama, $n=144$),

Name in paper	Checkpoint (Hugging Face ID)
Llama-3.1-8B	meta-llama/Llama-3.1-8B-Instruct
Qwen-2.5-7B	Qwen/Qwen2.5-7B-Instruct
Qwen3-8B	Qwen/Qwen3-8B
Gemma-2-9B	google/gemma-2-9b-it
Gemma-4-12B	google/gemma-4-12b-it
DeepSeek-R1-Dist.-7B	deepseek-ai/DeepSeek-R1-Distill-Qwen-7B
DeepSeek-R1-0528-8B	deepseek-ai/DeepSeek-R1-0528-Qwen3-8B
NLI encoder	FacebookAI/roberta-large-mnli
LLM judge	gpt-4o-mini (OpenAI API, temperature 0, August 2026)

Table 8: Exact model checkpoints. Benchmark traces originate from the FaithCoT-Bench release; the two open benchmark models are re-encoded with the matching checkpoints above.

judge and probe scores are essentially uncorrelated (Spearman 0.05), and the judge’s weakest per-model cell (0.599) sits below the probe’s out-of-fold 0.656 (recomputed for this paired analysis). A simple combination (mean of z-scored signals; the probe side uses out-of-fold predictions) outperforms the judge alone (0.699; paired bootstrap Δ [+0.02, +0.19], excluding zero, though the lower bound is close to it). For completeness: a cross-validated two-feature logistic scores 0.666, and the probe on the judge’s middle score tercile is 0.536 ($n=48$); the combination evidence is positive but modest.

Judge ablations (§4). Same protocol, system prompt, and message skeleton; only the withheld block changes. *Question-only* (CoT withheld): full 0.487 [0.456, 0.519], incorrect regime 0.467 [0.421, 0.512], correct regime 0.552 [0.503, 0.600], score vs. incorrectness 0.503 [0.473, 0.533]—at or near chance in every cell (the correct-regime interval grazes 0.5 under the trace bootstrap and includes it under a question-clustered one); 54% of its scores sit at the midpoint 50, and the differentiated remainder carries no usable signal. *CoT-only* (question and options withheld): full 0.693 [0.662, 0.725], incorrect regime 0.661 [0.614, 0.704], correct regime 0.769 [0.720, 0.817], score vs. incorrectness 0.560 [0.530, 0.590]. The blind-regime signal survives nearly intact without the question (0.661 of 0.679) while correctness tracking collapses toward chance (0.684 $\rightarrow$ 0.560). Caveat: traces often restate the question internally, so CoT-only is an upper bound on question independence.

Judge composition decomposition (§5). Incorrect-regime per-model judge AUROCs: Llama 0.599, Qwen-2.5 0.641, Gemini-2.5-Flash 0.642, GPT-4o-mini 0.657 (n -weighted

within-model mean 0.634 vs. pooled 0.679). Per-domain (incorrect / correct): AQuA 0.880 / 0.905, HLE-Bio 0.636 / 0.714, LogiQA 0.624 / 0.799, TruthfulQA 0.604 / 0.724. Within-domain degradation Δ s: TruthfulQA +0.119, LogiQA +0.175, AQuA +0.025; HLE-Bio is 29% of the incorrect regime but 1.5% of the correct regime.

J Pre-Registered Independent Test (BonaFide)

Protocol. Endpoints, population rule, effect margin, statistics and interpretation were committed before any value was computed against the target labels; the analysis stage was run once, after all inputs completed. The unit of analysis is a *response*, identified by (question id, generator, CoT); labels come from a response’s own rows and require an explicit whole-trace label (1,057 of 2,177 responses carry only step labels and are excluded, not treated as faithful). Judge protocol matches §4 exactly: same system prompt verbatim, temperature 0, JSON-forced, shown the clean question, the CoT and the final answer, and never the gold answer or the biased prompt. Deletion-family metrics use a free-form answer-logprob readout (no options exist in this data) and are pre-registered as restricted to the four generators that fit our hardware.

Outcomes. Incorrect-answer population $n=1,113$ (945 unfaithful / 168 faithful); the correct-answer cell holds 7 responses, all unfaithful, and supports no test. Note this is *not* an independent instance of the correctness entanglement of §5: 99.4% of BonaFide’s explicitly-labelled responses are incorrect-answer responses to begin with, so the conditional carries no information (lift 1.00, against 1.74 on FaithCoT). BonaFide therefore supplies no informative external check of this association; the




Figure 7: Per-cell (signal $\times$ model $\times$ domain) AUROC with bootstrap 95% CIs, by regime; NLI is available benchmark-wide (all four models), answer-tracing for the two open models. Cells with $n < 10$ or a single class are omitted; per-cell n is small, so individual CIs are wide. The incorrect-regime cells scatter around chance across models and domains, including the closed models; the correct-regime cells lean above chance across domains and in the closed models as well.

separate exploratory GRACE check is reported in Appendix K.

Signal	AUROC	95% CI
Judge (prompt A, primary)	0.419	[0.372, 0.464]
Judge (prompt B)	0.482	[0.437, 0.530]
Judge (gpt-4o)	0.270	[0.231, 0.310]
Step count	0.878	[0.852, 0.902]
NLI unsupported steps	0.826	[0.795, 0.855]
Answer-tracing (soft)	0.284	[0.171, 0.401]
Prefix instability	0.172	[0.101, 0.265]

Controls (post-unblinding, labelled as such).

BonaFide’s classes differ in composition: its faithful-incorrect class is entirely instruct-model, 95% SimpleQA (160/168) and $2.5\times$ shorter, while the unfaithful class is 43% SimpleQA, 28% HLE and 29% DDXPlus. Conditioning on generator, source task and absolute word band separates the two kinds of signal cleanly:

Strata	judge	steps	words	NLI	pairs
pooled	0.419	0.878	0.939	0.826	100%
+model, task	0.260	0.633	0.703	0.694	3.3%
+word band	0.193	0.521	0.547	0.603	1.4%

The feature signals decay toward chance as

comparisons are restricted, while the judge inversion *strengthens*—the asymmetry, not the pooled ranking, is the result. We therefore do not read step count as a working detector. Composition is not a complete explanation either: within SimpleQA alone step count is still 0.806. The judge inversion also holds instruct-only (0.332 [0.285, 0.381]; gpt-4o 0.184) and in four of five generators separately. Only 1.4% of cross-class pairs survive the strictest conditioning, so those estimates are correspondingly restricted. Stratified estimator and word bands follow the independent audit of our frozen predictions.

What is and is not established. Established: on this population the tested judges rank the native labels in the wrong direction, and that reversal survives the tested restrictions by generator identity, source task and length. Those restrictions retain limited pair support and do not identify its cause. Not established is any claim that nothing discriminates here: NLI unsupported-step count retains 0.603 [0.570, 0.638] under the strictest condition-

	Prev.	Ans.-tracing (inv.)	NLI
Incorrect rgm.	0.544	0.565 [0.489, 0.650]	0.586 [0.510, 0.677]
Correct rgm.	0.204	0.367 [0.271, 0.488]	0.350 [0.253, 0.448]

Table 9: AUPRC robustness for the Table 5 comparisons (prefix instability behaves like answer-tracing: 0.518/0.385). Random baseline = prevalence. In the incorrect regime AUPRC matches prevalence (consistent with the null); in the correct regime it reaches 1.7–1.9× prevalence, so the moderate AUROCs are not masking failed minority-class retrieval.

ing. Not established: *why*. Judged scores are higher for short honest errors (mean 91/100 under gpt-4o) than for elaborate unacknowledged-hint rationalizations (61/100), which is consistent with several explanations, including a rubric that differs from the native process/source-claim criteria. Reliance or nondisclosure alone does not reproduce the full native target. Independent human checking of alleged violations remains outstanding. The exploratory cached rubric/evidence comparison below does not isolate a sufficient explanation. All tested judges are from one provider.

Exploratory reconciliation on common responses. We compare cached A/B/C scores and historical GPT-4o scores D on exactly 150 common responses (123 questions; 85 unfaithful, 65 faithful). A uses a generic rubric with the full prompt; B uses that rubric with the local hint-stripping rule; C uses our rubric with the full prompt; D uses our rubric with the clean question. B is not an exact implementation of the upstream definition-aware no-hint monitor. All score directions are fixed towards unfaithfulness.

Arm	AUROC	95% interval
A	0.319	[+0.229, +0.412]
B	0.400	[+0.307, +0.490]
C	0.367	[+0.278, +0.456]
D	0.331	[+0.242, +0.418]

Direct paired differences are A-B: -0.081 , $[-0.178, +0.015]$; A-C: -0.048 , $[-0.128, +0.036]$; C-D: $+0.036$, $[-0.023, +0.093]$. These intervals do not establish a context or rubric remedy, nor equivalence. A on its larger, unmatched sample is descriptive only. On the 143 common incorrect answers, A/B/C/D AUROCs are 0.335/0.419/0.370/0.332. B minus D is $+0.087$ ($[+0.006, +0.161]$), an exploratory nominal interval among multiple contrasts involving a historical comparator; this is not independent evidence of a general correction method. The reconciliation exposes 204 of the original 807 evaluation responses across 156 question clusters; 397 evaluation responses share

those clusters. We retain that split and flag exposure. The seven correct answers lie outside the split’s incorrect-answer population. Legacy scores lack raw judge outputs and exact request provenance, so these analyses cannot validate quoted rationales or attribute differences solely to evidence access.

K Second External Standard (GRACE)

GRACE (Pham et al., 2026) supplies human step-level faithfulness labels on 437 released test traces (2,044 steps; 404 source questions) over four context-grounded tasks. The selected verifier-disagreement population is not a natural prevalence sample, and prior exposure to the preliminary release precludes treating this as an untouched confirmatory test. Unlike BonaFide’s seven correct, all-unfaithful responses, GRACE permits an exploratory correctness-conditioned analysis.

Our primary endpoint is Spearman correlation with the fraction of native unfaithful steps. Five binary aggregation rules (any, more than one quarter, at least half, more than half, and all) are sensitivity endpoints. We inspected aggregation counts before analysis and claim no preregistration. Correlations and direct correctness-stratum differences use 2,000 question-cluster bootstrap resamples, grouping by dataset and original question ID. Directions are fixed; mean entailment and citation fraction are negated once.

The cached RoBERTa NLI protocol limits each premise to 4,000 characters and each premise/step token pair to 512 tokens. The context arm uses supplied passages; the prior-step arm uses the question and preceding steps, omitting the passages. This is a cross-check of these limited NLI procedures, not the complete detector panel. Original request and token-truncation logs are unavailable for this legacy cache. No new GRACE judge inference is represented by these results.

Correctness and coverage. We use conservative exact-answer and explicit-option checks, followed by recorded assistant answer-equivalence review for unresolved responses. This is exploratory review, not independent human gold. Of 437 traces, 260 are resolved automatically; review yields 280 correct and 154 incorrect answers, leaving 3 unresolved. The checker does not accept first-character matches, answer substrings, or token-overlap scores as correctness labels. Refusals and contradictory commitments are inspected explicitly. Reference answers themselves have not been independently revalidated.

Correctness-conditioned results.

Signal (ρ)	Correct	Incorrect	Difference
Context NLI, unsup. rate	+0.241	+0.216	+0.025
Context NLI, neg. entail.	+0.273	+0.159	+0.113
Prior NLI, unsup. rate	+0.081	+0.085	-0.004

Paired question-bootstrap 95% intervals for these differences, respectively, are $[-0.177, +0.219]$, $[-0.088, +0.312]$, $[-0.220, +0.210]$. These comparisons do not resolve a correctness-stratum difference for the tested NLI procedures. The intervals also allow meaningful differences; they do not establish equivalence or a mechanism unique to FaithCoT. On all 437 traces, independently of correctness checking, context unsupported-rate and negative mean-entailment correlations are +0.226, +0.254. These are pooled grounding associations, not evidence that every detector family transfers.

Sensitivity and limits. The automatic-only subset has 197 correct and 63 incorrect answers; its three regime differences are -0.131 , -0.141 , -0.046 . This subset changes task and answer-form composition substantially and is not an independent replication. We also release task-specific fraction/correlation analyses and every assignment of the unresolved answers; assignment sensitivity does not bound mistakes in resolved reviews. The descriptive lift $P(\text{incorrect} \mid \text{unfaithful})/P(\text{incorrect})$ ranges from 1.16 to 1.71 over the five rules; at at-least-half unfaithful it is 1.61 (95% interval $[+1.432, +1.804]$). The all-steps-unfaithful rule gives an interval $[+0.832, +1.723]$. Thus the association depends on the aggregation and selected population; evidence for it is not uniformly resolved across rules. Length remains associated with the fraction target (step-count $\rho = -0.223$); step-count AUROC ranges from 0.184 to 0.491 across aggregations. Dataset-by-track conditional pair AUROCs

are provided as a separate binary endpoint with pair coverage; they are not adjusted Spearman correlations.

The earlier 40-example sample analysis is historical development evidence and is superseded here by the released test-set analysis; it is not an additional independent replication.